

Large-Scale and Multi-Structured Databases

MongoDB Java Driver

Prof. Pietro Ducange

Ing. Alessio Schiavo

alessio.schiavo@phd.unipi.it


Copyright Issues

Most of the information included this presentation have been extracted from the official documentation of MongoDB Java Driver (<http://mongodb.github.io/mongo-java-driver/>).

MongoDB Java Driver

This driver allows us to manipulate using Java Application data stored in a MongoDB database.

The main suggestion is to follow the official documentation available at <https://www.mongodb.com/docs/drivers/java/sync/current/>



Introduction

The official MongoDB Java Drivers providing both synchronous and asynchronous interaction with MongoDB.

Features

BSON Library

A standalone BSON library, with a new Codec infrastructure that you can use to build high-performance encoders and decoders without requiring an intermediate Map instance.

MongoDB Driver

An updated Java driver that includes the legacy API as well as a new generic MongoClient interface that complies with a new cross-driver CRUD specification.

MongoDB Reactive Streams Driver

Providing asynchronous stream processing with non-blocking back pressure for MongoDB. Fully implements the [Reactive Streams API](#) for providing interop with other reactive streams within the JVM ecosystem.

Releases

RELEASE	DOCUMENTATION	
4.8.0-beta0	Reference	API
4.7.2	Reference	API
4.6.0	Reference	API
4.5.1	Reference	API
4.4.1	Reference	API
4.3.4	Reference	API
4.2.3	Reference	API
4.1.2	Reference	API

Installation

- The recommended way to get started using *one of the drivers* in your project is with a dependency management system, such as **MAVEN**.
- The current official indication is to use the ***mongodb-driver-sync*** in the Java application.
- Specify in the pom file the following dependency

```
<dependencies>
....
<dependency>
  <groupId>org.mongodb</groupId>
  <artifactId>mongodb-driver-sync</artifactId>
  <version>4.7.2</version>
</dependency>
...
</dependencies>
```


Making a Connection

- The ***MongoClients.create()***, allows us to make a connection to a running MongoDB instance.
- A simple way to connect to a MongoDB server is:

`MongoClients.create(connectionString)`



```
MongoClient myClient = MongoClient.create("mongodb://localhost:27017");
```



Some Important Considerations

- The ***MongoClient instance*** represents a ***pool of connections*** to the database; we will only need one instance of class MongoClient even with ***multiple threads***.
- **IMPORTANT:** MongoClient is thread-safe. Multiple access to the single instance is managed by the class itself
- Typically, for simple projects, we only create ***one MongoClient instance*** for a given MongoDB deployment (e.g. standalone, replica set, or a sharded cluster) and use it ***across the whole application***.
- Call ***MongoClient.close()*** to clean up resources at the end of the application.

Access a Database

- Once we have a MongoClient instance connected to a MongoDB deployment, we can use the ***MongoClient.getDatabase()*** method to access a database.
- ***Specify the name*** of the database to the getDatabase() method. If a database does not exist, MongoDB creates the database when you first store data for that database.
- The following example accesses the mydb database:

```
MongoDatabase database = mongoClient.getDatabase("mydb");
```

Access a Collection

- Once we have a ***MongoDatabase*** instance, we can use its ***getCollection()*** method to access a collection.
- ***Specify the name*** of the collection to the `getCollection()` method. If a collection does not exist, MongoDB creates the collection when you first store data for that collection.
- For example, using the database instance, the following statement accesses the ***collection*** named ***test*** in the mydb database:

```
MongoCollection<Document> collection = database.getCollection("test");
```


Code overview

```
import com.mongodb.client.*;
import com.mongodb.ConnectionString;

//Create connection string
ConnectionString uri = new ConnectionString("mongodb://localhost:27017");
//Create a mongoDB client
MongoClient myClient = MongoClient.create(uri);

//Connect to mydb database
MongoDatabase database = mongoClient.getDatabase("mydb");

//Select the collection test
MongoCollection<Document> collection = database.getCollection("test");

...
//insert, remove, update elements to/from the collection
...

//Close mongoDB connection and release resources
myClient.close();
```

Handling Collections basics

- To **create a Collection using the Java driver**, we can use the ***createCollection*** method of a *MongoDatabase* instance. For example, let us create a collection called “exampleCollection”:

```
database.createCollection("exampleCollection");
```

- To **get list of existing Collections using the Java driver**, we can use the *MongoDatabase.listCollectionNames()* method:

```
for (String name : database.listCollectionNames()) {  
    System.out.println(name);  
}
```

- To **drop a Collection using the Java driver**, we can use the *MongoCollection.drop()* method:

```
MongoCollection<Document> collection = database.getCollection("exampleCollection");  
collection.drop();
```

Create a Document

- To create the document using the Java driver, we can use the **Document class**. For example, consider the following JSON document:

<i>string</i>	{ "name" : "MongoDB",
<i>number</i>	"count" : 1,
<i>list</i>	"versions": ["v3.2", "v3.0", "v2.6"],
<i>embedded obj</i>	"info" : { x : 203, y : 102 } }

```
Document doc = new Document("name", "MongoDB")
    .append("count", 1)
    .append("versions", Arrays.asList("v3.2", "v3.0", "v2.6"))
    .append("info", new Document("x", 203).append("y", 102));
```

- Otherwise, you can use a string that represent the json file (Be careful to escape double quotes!)

```
Document doc = Document.parse("{name: \"Alessio\", surname: \"Schiavo\"}");
```

Insert Document

- To insert a single document into the collection, we can use the collection's `insertOne()` method.

```
collection.insertOne(doc);
```

- To insert a set of documents, contained into a list of documents we can use the following example of code:

```
List<Document> documents = new ArrayList<Document>();
```

```
[...populate documents...]
```

```
//Insert multiple documents
```

```
collection.insertMany(documents);
```

```
//count the # of docs in a collection
```

```
System.out.println(collection.countDocuments());
```

Query a Collection

- To query a collection, we can use the collection's ***find() method***.
- We can call the method without any arguments ***to query all documents*** in a collection or ***pass a filter*** to query for documents that ***match the filter criteria***.
- The following example retrieves all documents in the collection and prints the returned documents:

```
try (MongoCursor<Document> cursor = myColl.find().iterator())  
{  
    while (cursor.hasNext())  
    {  
        System.out.println(cursor.next().toJson());  
    }  
}
```

Show results of a query

- We can iterate through query results by using a consumer function (statically or locally defined)

```
import java.util.function.Consumer;

Consumer<Document> printDocuments = doc ->
    {System.out.println(doc.toJson());};

myColl.find().forEach(printDocuments);
```

- Collect results in a list

```
List<Document> results =
    myColl.find().into(new ArrayList<>);
```

Specify a Query Filter

- To query for documents that match certain conditions, pass a ***filter object*** to the `find()` method.
- To facilitate creating filter objects, Java driver provides the Filters helper ([link](#))
- The following example retrieves all documents in the collection where $50 < i \leq 100$ and prints the returned documents:

```
try (MongoCursor<Document> cursor =  
    myColl.find(and(gt("i", 50), lte("i", 100))).iterator())  
{  
    while (cursor.hasNext())  
    {  
        System.out.println(cursor.next().toJson());  
    }  
}
```

Update Documents

- To update documents in a collection, we can use the collection's ***updateOne*** and ***updateMany*** methods.
- These functions needs two parameters:
 - ***A filter object*** to determine the document or documents to update.
 - ***An update document*** that specifies the modifications. Check the [manual](#) for a list of the available operators

```
collection.updateOne(  
    eq("name", "Alessio"),  
    set("age", 28));
```


Update Documents: Examples

- The following example updates the first document that meets the filter *i equals 10* and **sets** the value of *i to 110*:

```
collection.updateOne(eq("i", 10), set("i", 110));
```

- The following example **increments the value of i by 100** for all documents where the value of field *i is less than 100*:

```
UpdateResult updateResult =  
    collection.updateMany(lt("i", 100), inc("i", 100));  
System.out.println(updateResult.getModifiedCount());
```

- The update methods return an [UpdateResult](#) which provides information about the operation including the number of documents modified by the update.

Delete Documents

- To delete documents from a collection, we can use the collection's ***deleteOne*** and ***deleteMany*** methods.
- As a parameter it requires just a ***filter object*** to select the documents to delete. The example deletes at ***most one document*** that meets the filter ***i equals 110***:

```
collection.deleteOne(eq("i", 110));
```

- The following example deletes ***all documents*** where ***i is greater or equal to 100***:

```
DeleteResult deleteResult = collection.deleteMany(gte("i", 100));  
System.out.println(deleteResult.getDeletedCount());
```

- The delete methods return a **DeleteResult** which provides information about the operation including the number of documents deleted.

Aggregation pipeline

- To perform aggregation, pass a list of aggregation stages to the ***MongoCollection.aggregate()*** method.
- For a complete list of aggregations, check the reference [documentation](#)

```
Bson myMatch = Aggregates.match(Filters.eq("categories", "Bakery"));  
Bson myGroup = Aggregates.group("$stars", Accumulators.sum("count", 1));  
  
collection.aggregate(Arrays.asList(myMatch, myGroup))  
    .forEach(printDocuments());
```

Aggregation pipeline (2)

- Import static filters, aggregations, projections and accumulators to improve readability of your code

```
Import static com.mongodb.client.model.Aggregates.*;
Import static com.mongodb.client.model.Accumulators.*;
Import static com.mongodb.client.model.Projections.*;
Import static com.mongodb.client.model.Filters.*;

Bson myMatch = match(eq("categories", "Bakery"));
Bson myGroup = group("$stars", sum("count", 1));

collection.aggregate(Arrays.asList(myMatch, myGroup))
    .forEach(printDocuments());
```

Pipeline - match

- The **match** operator filters the documents to pass only the ones that match the specified condition(s) to the next pipeline stage.

```
//Strings
Bson myMatch = match(eq("categories", "Bakery"));

//Integers
Bson myMatch2 = match(gte("pop", 50000));

//Logic operators
Bson myMatch3 = match(
    and(eq("name", "XYZ Coffee Bar"), eq("categories", "Coffee")));

collection.aggregate(Arrays.asList(myMatch, myMatch2, myMatch3))
    .forEach(doc -> System.out.println(doc.toJson()));
```

Pipeline - group

- The **group** operator groups input documents by the specified `_id` expression (first argument) and for each distinct grouping. It returns a document.
- This operator can include accumulators (sum, avg, max, min, ...)

```
Bson groupSingle = group("$city", sum("totPop", "$pop"));
```

- For **multiple fields grouping** it is better to define directly a document:

```
Bson groupMultiple = new Document("$group",  
    new Document("_id", new Document("city", "$city")  
        .append("state", "$state"))  
    .append("totalPop", new Document("$sum", "$pop")));
```

```
//Same as defining the following document
```

```
//{$group: {_id:{city:$city, state:$state}, totalPop:{$sum:$pop}}}
```

Pipeline - project

- The **project** operator passes along the documents with the requested fields to the next stage in the pipeline.
- The specified fields can be existing fields from the input documents or newly computed fields.
- I can exclude, include or compute new fields

```
Bson p1 = project(fields(include("city")));  
Bson p2 = project(fields(  
    excludeId(), include("city", "state")));  
Bson p3 = project(fields(  
    include("city"), exclude("state")));  
Bson p4 = project(fields(  
    computed("myID", "_id"), include("city")));
```

Pipeline - sort

- The **sort** operator Sorts all input documents and returns them to the pipeline in sorted order.
- The order can be *ascending* or *descending*

```
Bson s1 = sort(descending("pop"));  
Bson s2 = sort(ascending("city"), descending("pop"));  
  
collection.aggregate(Arrays.asList(s2))  
    .forEach(printDocuments());
```


Pipeline – limit, skip

- The **limit** operator limits the number of documents passed to the next stage
- In conjunction with **limit**, we can implement queries that search, for example, for the top 3 biggest cities
- The **skip** operator, instead, skips over the specified number of documents that pass into the next stage

```
//Find the 2nd and 3rd biggest cities
```

```
Bson mySort = sort(descending("pop"));
```

```
Bson myLimit = limit(2);
```

```
Bson mySkip = skip(1);
```

```
collection.aggregate(Arrays.asList(mySort, mySkip, myLimit))  
    .forEach(printDocuments());
```

Pipeline – unwind

- The **unwind** operator deconstructs an array field from the input documents to output a document for *each* element
- Each output document identical to the input doc except for the value of the *grades* field which now holds a value from the original *grades* array:

```
//Find the average score for the American cuisine
Bson m = match(eq("cuisine", "American"));
Bson u = unwind("grades");
Bson g = group("$cuisine", avg("avgGrade", "$grades"));

collection.aggregate(Arrays.asList(m, u, g))
    .forEach(printDocuments());
```

Create Indexes

- To create an index on a field or fields, pass an index specification document to the ***createIndex()*** method.
- An index key specification document contains the ***fields*** to index and the ***index type*** for each field:

new Document(<field1>, <type1>).append(<field2>, <type2>) ...

- For an **ascending** index type, specify **+1** for <type>. For a **descending** index type, specify **-1** for <type>.
- The following example creates an ascending index on the i field:

```
collection.createIndex(new Document("city", 1));
```

Exercises

ZIPs dataset

1. Find zip codes of Texan cities.
2. Find cities' zips with a population of at least 100'000, but no more than 200'000.
3. Find the 5 most populated cities.
4. For each state, find the average population of its cities.

POSTS dataset

1. Find posts published after 2012-11-20.
2. Find posts with the tag 'computer'.
3. Find, for each tag, the total number of posts.
4. Find the top three commentators according to the number of comments.
5. Find the most versatile commentator. *Versatile* means that he/she commented on the highest number of *distinct* topics (a.k.a. tags). For a tag to count, he/she must have at least five comments about that topic.

Help: You can find useful examples on ZIPS and POSTS datasets [here](#).

Suggested Readings

Students are invited to read the official documentation of MongoDB.

The documentation is available at (latest version 4.1.1):

<http://mongodb.github.io/mongo-java-driver/>

Check here <https://docs.mongodb.com/ecosystem/drivers/> for drivers for ***different programming languages*** and their details.