


Petri nets

- **Petri nets (PN)**
high-level modelling formalism for concurrent and distributed systems (Petri, Carl A. (Ph. D. Thesis), University of Bonn, 1962)
- **Stochastic Petri nets (SPN)**
temporal and probabilistic information added to the model
the approach aimed at equivalence between SPN and Markov Chains
idea of associating an exponentially distributed random delay with the PN transitions (1990)
- **Stochastic Activity networks (SAN)**
extension to Stochastic Petri nets (data structure, complex firing rules, global variables,) (2001)

Petri nets

A Petri net consists of places, transitions, and arcs. Arcs are from a place to a transition or vice versa. Places may contain a discrete number of marks called tokens.


Graphically:



- the places with an arc that runs from the place to a transition are called the input places of the transition
- the places with an arc that runs from the transition to a place are called the output places of the transition

Petri nets

Marking of the network:
any distribution of tokens over the places. This
represents a configuration of the network.



A transition of a Petri net may fire if it is enabled, i.e., there are sufficient tokens in all of its input places
This is named the weight of the arc (by default 1).

When the transition fires, it consumes the required input tokens, and creates tokens in its output places.

Petri nets are well suited for modeling the concurrent behavior of distributed systems.

Petri nets

$$N_{P/T} = \langle P, T; F, W, M_0 \rangle$$

$$W : F \rightarrow \mathbf{IN} - \{0\}$$

$$M_0 : P \rightarrow \mathbf{IN}$$

P: places

T: transitions

F: arcs

W: weight

M_0 : initial marking


t, y transitions

$$\bullet t = \{p \in P \mid \langle p, t \rangle \in F\} \quad \text{preset}$$

$$y^\bullet = \{z \in P \mid \langle y, z \rangle \in F\} \quad \text{postset}$$

$$t \text{ enabled if: } \forall_{p \in \bullet t} M(p) \geq W(\langle p, t \rangle)$$

$$\text{we write } M[t\rangle$$



Transition firing

Firing rule

$$\forall_{p \in (\bullet t - t \bullet)} \quad M'(p) = M(p) - W(\langle p, t \rangle)$$


$$\forall_{p \in (t \bullet - \bullet t)} \quad M'(p) = M(p) + W(\langle t, p \rangle)$$

$$\forall_{p \in (\bullet t \cap t \bullet)} \quad M'(p) = M(p) - W(\langle p, t \rangle) + W(\langle t, p \rangle)$$

$$\forall_{p \in P - (\bullet t \cup t \bullet)} \quad M'(p) = M(p)$$

We write $M [t\rangle M'$

$M_0 [t\rangle M_1$



$M_0 = (2, 3, 0, 0)$

$M_1 = (1, 0, 2, 1)$

The firing of a transition is atomic
(a single non-interruptible step)


Transition sequence

Reachable marking $R_N(M)$:

$$M \in R_N(M)$$

$$M' \in R_N(M) \wedge \exists_{t \in T} M' [t \rangle M'' \Rightarrow M'' \in R_N(M)$$

Reachability graph



Transition sequence

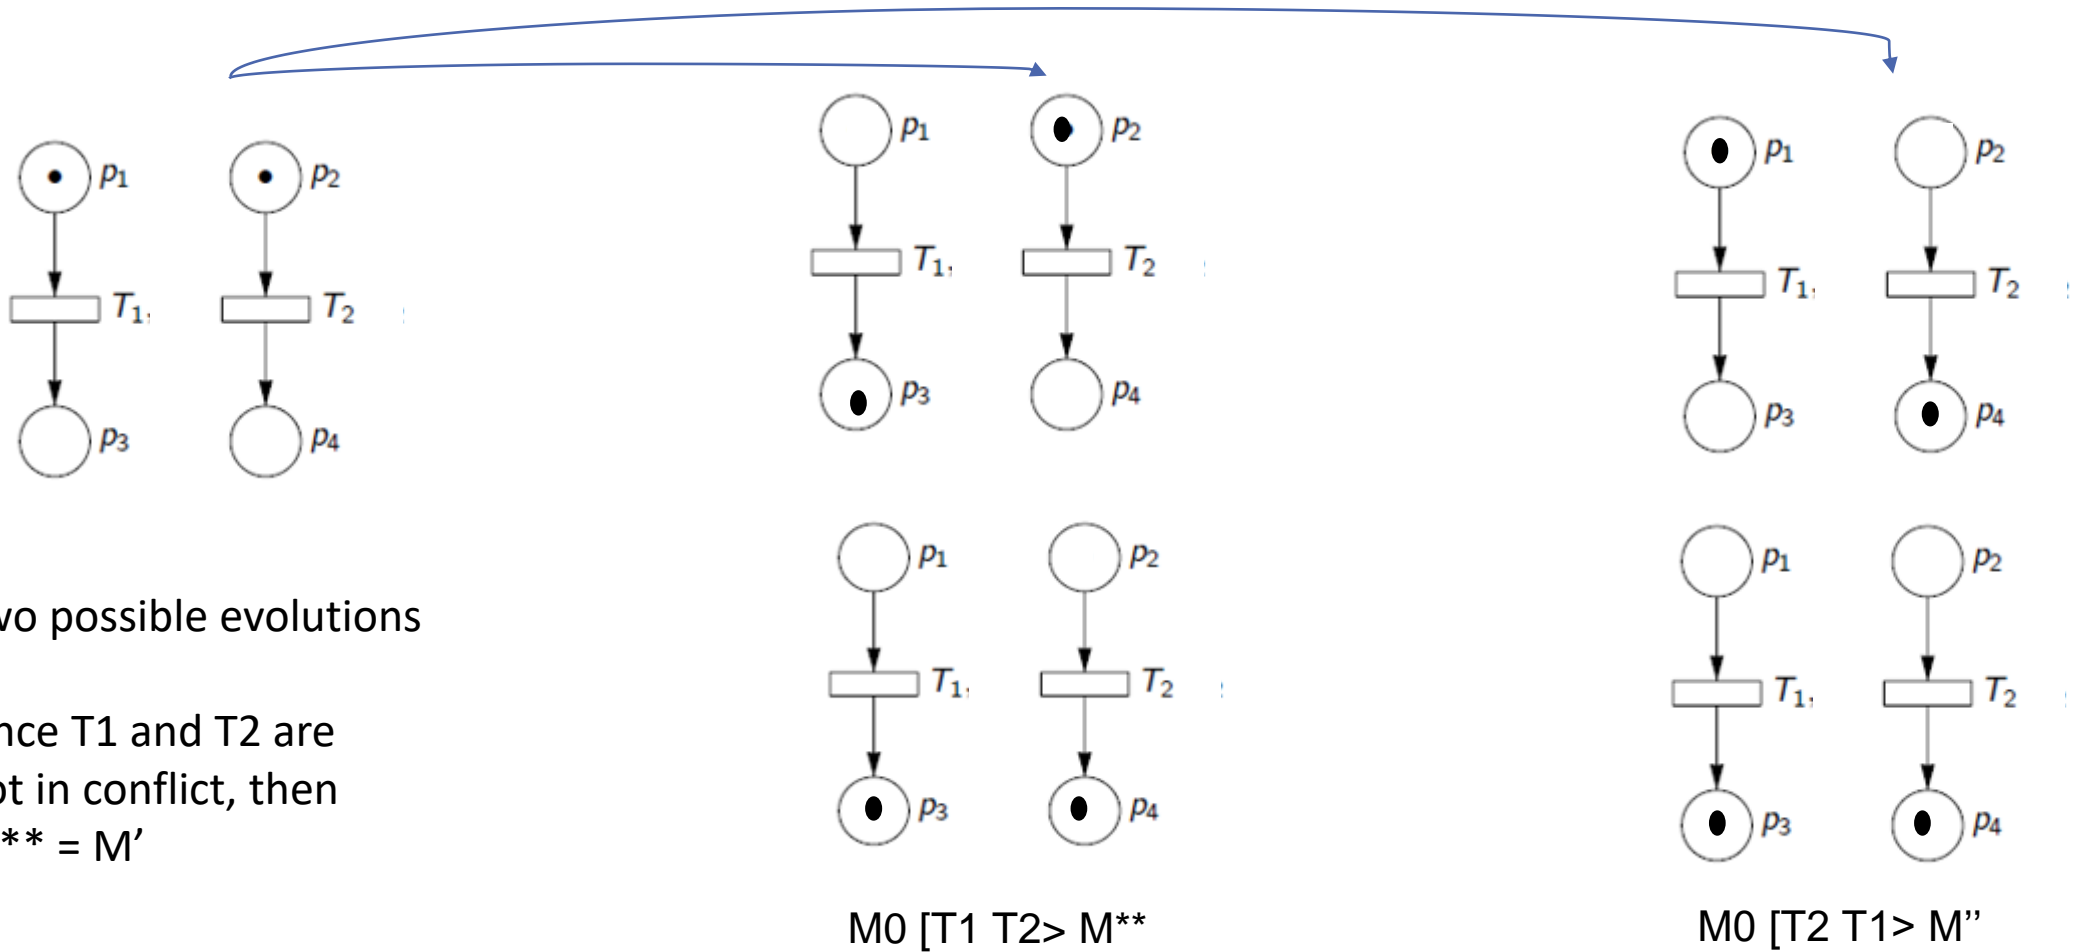
Let $M [t_1 \rangle M' \wedge M' [t_2 \rangle M''$ then $M [t_1 t_2 \rangle M''$

If $M [t_1 \dots t_n \rangle M^{(n)}$ then $t_1 \dots t_n$ is a transition sequence

Evolution of the net



At any time, one of the enabled transitions is executed at a time

If multiple transitions are enabled, these transitions will fire in any order. The execution of Petri nets is nondeterministic



Conflict resolution


T1, T2 are in conflict
For example, when T1 fires,
T2 is no more enabled



$M0=(1,0,0,)$ $M1=(0,1,0)$

$M0 [T1 > M1$

T2 not enabled in M1




$M0=(1,0,0,)$ $M2=(0,0,1)$

$M0 [T2 > M2$

T1 not enabled in M2

Non-sequential processes




Building a process step by step



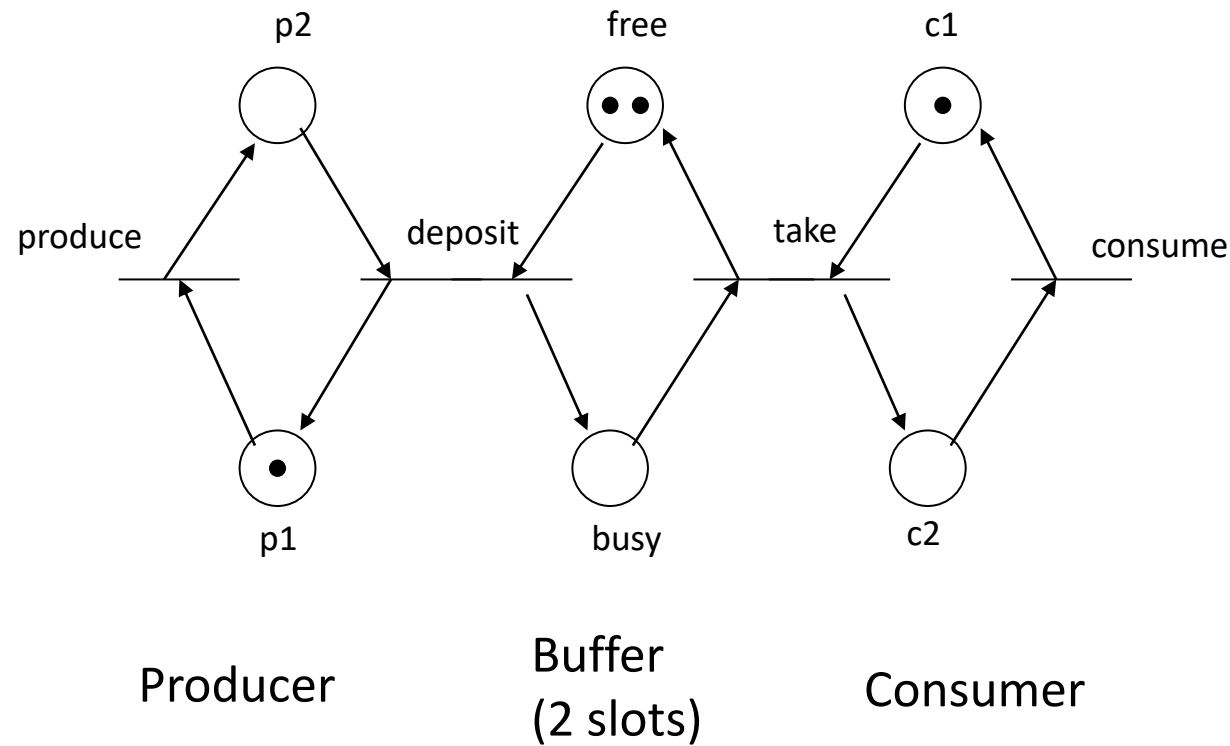
Petri net

Simulation of concurrent and distributed systems

degree of parallelism: maximum number of transitions enabled in a cut of the process

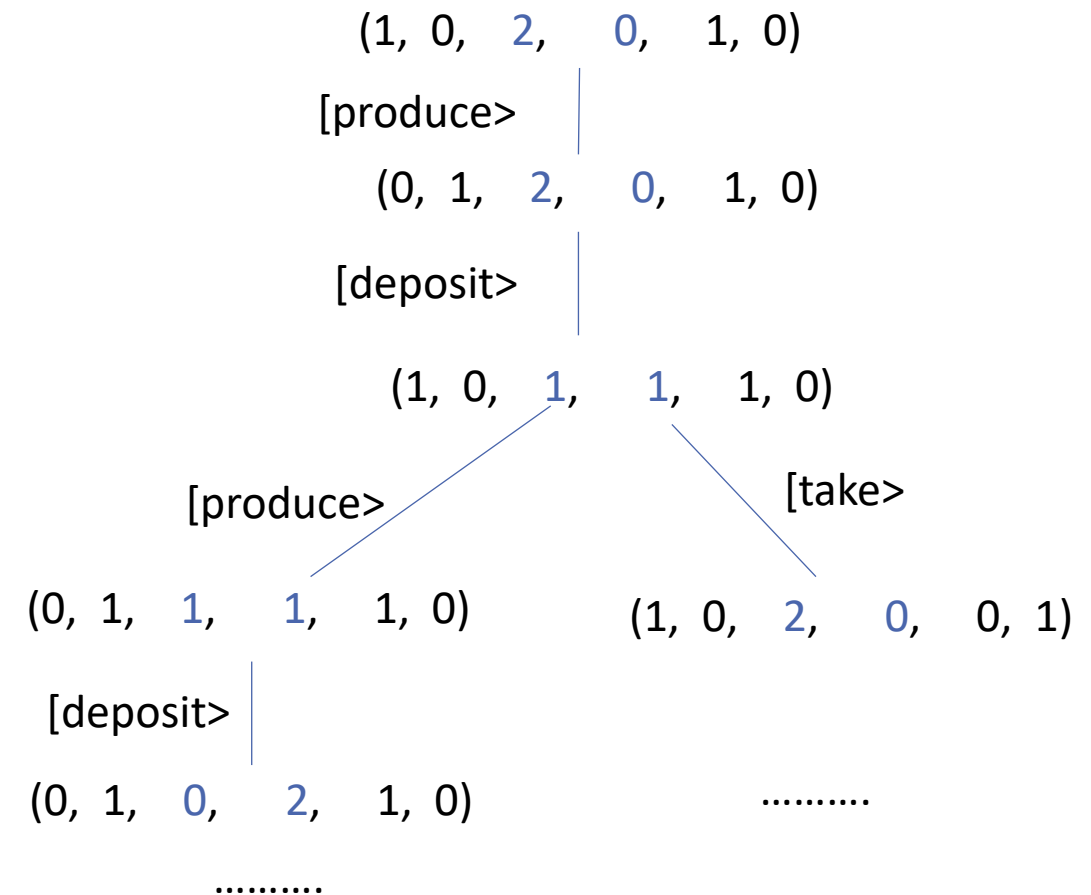


Producer/Consumer example



$$\text{free} + \text{busy} = 2$$

$M = (p1, p2, \text{free}, \text{busy}, c1, c2)$



Analysis

- Reachability graph
the reachability graph may be infinite, i.e. if there is no bound on the number tokens on some place.
The coverability graph computes an over-approximation (remains exact as long as the number of markings is finite)
- Conditions on reachable markings
used to check whether some 'bad' state can occur (e.g., violation of mutual exclusion property)
- Transition invariants
indicate a possible loop in the net (a sequence of transitions which leads back to the marking it starts in)
- Place invariants
global properties of markings always satisfied: the number of tokens in all reachable markings satisfies some linear invariant (e.g., all reachable markings M satisfy: $M(p1) + M(p4) = 5$)
- Deadlock
A marking is called a deadlock if no transition is enabled in the marking. A net is deadlock-free if no reachable marking is a deadlock
-

Dependability evaluation and Petri nets

System description with Petri nets

- Place:
 - a system component (one for every component)
 - a class of system components (CPU , Memory, ..)
 - components in a given state (CPU, FaultyCPU, ..)
 -
- Token:
 - number of components (number of CPUs)
 - Occurrence of an event (fault, ..)
 -
- Transitions:
 - Occurrence of an event (Repair, CPUFaulty, ...)
 - Execution of a computation step
 - ...

Timed transitions

- Timed transition: an activity that needs some time to be executed
 - assigning a delay (**local time d**) to each transition
 - assigning a **global time** to the PN
- When a transition is enabled a local timer is set to **d**
 - the timer is decreased
 - when the timer elapses, the transition fires and remove tokens from input places
- If the transitions is disabled before the firing, the timer stops

Handling of the **timer** (two alternatives):

Continue:

the timer maintains the value and it will be used when the transition is enabled again

Restart:

the timer is reset

Timed transitions

Sequence of timed transitions:

$$(\tau_{k1}, t_{k1}) \dots (\tau_{kn}, t_{kn})$$

where

- $\tau_{k1} \leq \tau_{k2} \leq \tau_{kn}$
- $[\tau_{ki}, \tau_{ki+1})$

is the period of time between the firing of two transitions

in this period of time the net does not change the marking

STOCHASTIC PETRI NET:

when the **delay d** of a timed transition is a **random variable**


Stochastic Petri nets (SPN)

when the delay is a negative exponential distribution random variable,
the reachability graph is a Markov chain

Transition rate: λ_k is the transition rate (positive real number) of transition T_k

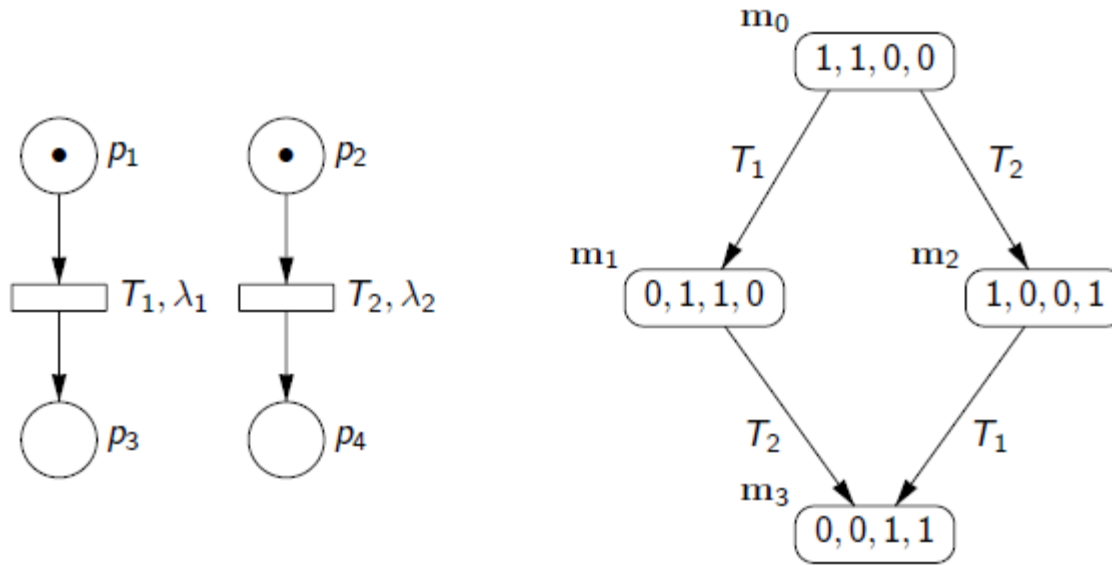
- delay of transition T_1 random variable exponentially distributed with parameter λ_1
- delay of transition T_2 random variable exponentially distributed with parameter λ_2

Assumption: delay of transitions independent



Stochastic Petri nets (SPN)


Reachability graph



a timed transition T enabled at time t , with d the random value for the transition delay, fires at time $t+d$, if it remains enabled in the interval $[t, t+d)$

Stochastic Petri nets (SPN)

Conflict resolution





the resolution of the conflict depends on the delay of T_1 and T_2

A redundant system with repair

- Two identical CPUs
- Failure of the CPU: exponentially distributed with parameter λ
- Fault detection: exponentially distributed with parameter δ
- CPU repair: exponentially distributed with parameter μ

SPN



Reachability graph


A redundant system with repair

Markov chain

failure rate = λ

detection rate = δ

repair rate = μ




Properties

- Steady-state probability that both processors behave correctly
- Steady-state probability of one undetected faulty processor
- Steady state probability that both processors must be repaired
-

M. Ajmone Marsan. Stochastic Petri nets: An elementary introduction. In G. Rozenberg, editor, *Advances in Petri Nets 1989*, volume 424 of *LNCS*, pages 1–29. Springer Verlag, 1990.

Example


Multiprocessor system with 2 processors and 3 shared memories system.
System is operational if at least one processor and one memory are operational.



λ_m failure rate for memory
 λ_p failure rate for processor

μ_m repair rate for memory
 μ_p repair rate for processor

Example




System state

$M = (M_{work}, P_{work}, M_{faulty}, P_{faulty})$

Initial marking $M_0 = (3, 2, 0, 0)$

$(3, 2, 0, 0) [T_{fp}]> (3, 1, 0, 1) [T_{fm}]> (2, 1, 1, 1) \dots\dots$

Example



Stochastic Activity Networks (SAN)


A variant of stochastic Petri nets

A system is described with SAN through four disjoint sets of nodes:

- places
- input gates
- output gates
- activities

Cases can be used to model probabilistic behaviors.

Graphically, places are drawn as circles, input (output) gates as left-pointing (right-pointing) triangles, instantaneous activities as narrow vertical bars, and timed activities as thick vertical bars. Cases are drawn as small circles on the right side of activities.




W. H. Sanders and J. F. Meyer. Stochastic activity networks: Formal definitions and concepts. In E. Brinksma, H. Hermanns, and J. P. Katoen, editors, *Lectures on Formal Methods and Performance Analysis*, volume 2090 of *LNCS*, pages 315–343. Springer Verlag, 2001.

Stochastic Activity Networks (SAN)

activity:

- **instantaneous**
system activities that are completed in a negligible amount of time
- **timed** activities of the modeled system whose durations impact the system's ability to perform
(the duration is expressed via a time distribution function)



input gate:

- has a finite set of inputs and one output
- enabling predicate (boolean function on the marking)
- input function defined for all values for which the enabling predicate is true
(marking of input places update)

output gate:


has a finite set of outputs and one input
output function (marking update)

Gates are introduced to permit greater flexibility in defining enabling and completion rules
Each input or output gate is connected to a single activity

Stochastic Activity Networks (SAN)

Cases associated with activities permit the realization of two types of *uncertainty*:

- uncertainty about which activities are enabled in a certain state is realized by cases associated with instantaneous activities
- uncertainty about the next state assumed upon completion of a timed activity is realized by cases associated with that activity



cases and case distribution:

instantaneous or timed activity may have mutually exclusive outcomes, called cases, chosen probabilistically according to the case distribution of the activity. Output gates allow the definition of different next marking rules for the cases of the activity

Stochastic Activity Networks (SAN)

Other extensions to PN:

Modules


(describe a system as a composition/replication of modules)

Shared variables


(global objects that can be used to exchange information among modules)

Extended places

(places whose marking is a complex data structure instead of a non-negative integer)



Example




Gate	Definition
IG0	Enabling predicate if $((p1 \rightarrow \text{Mark}() > 0 \wedge p2 \rightarrow \text{Mark}() == 0) \parallel$ $(p1 \rightarrow \text{Mark}() == 0 \wedge p2 \rightarrow \text{Mark}() > 0))$ return 1; else return 0;
OG0	Output function $p4 \rightarrow \text{Mark}()++;$ if $(p3 \rightarrow \text{Mark}() == 0)$ $p3 \rightarrow \text{Mark}()=1;$

Enabling predicates, and input and output gate functions are usually expressed using pseudo-C code.

Example

Example of Case: *activity generate_input*
Bit 0 or bit 1 is generated according to the probability assigned to the cases of the activity



The behavior of an activity network is a characterization of the possible completions of activities, selection of cases, and changes in markings.

Example

SAN evolution starting from a given marking M

- (i) the instantaneous activities enabled in M complete in some unspecified order;
- (ii) if no instantaneous activities are enabled in M , the enabled (timed) activities become active;
- (iii) the completion times of each active (timed) activity are computed stochastically, according to the respective time distributions; the activity with the earliest completion time is selected for completion;
- (iv) when an activity (timed or not) completes, the next marking M' is computed by evaluating the input and output functions;
- (v) if an activity that was active in M is no longer enabled in M' , it is removed from the set of active activities.

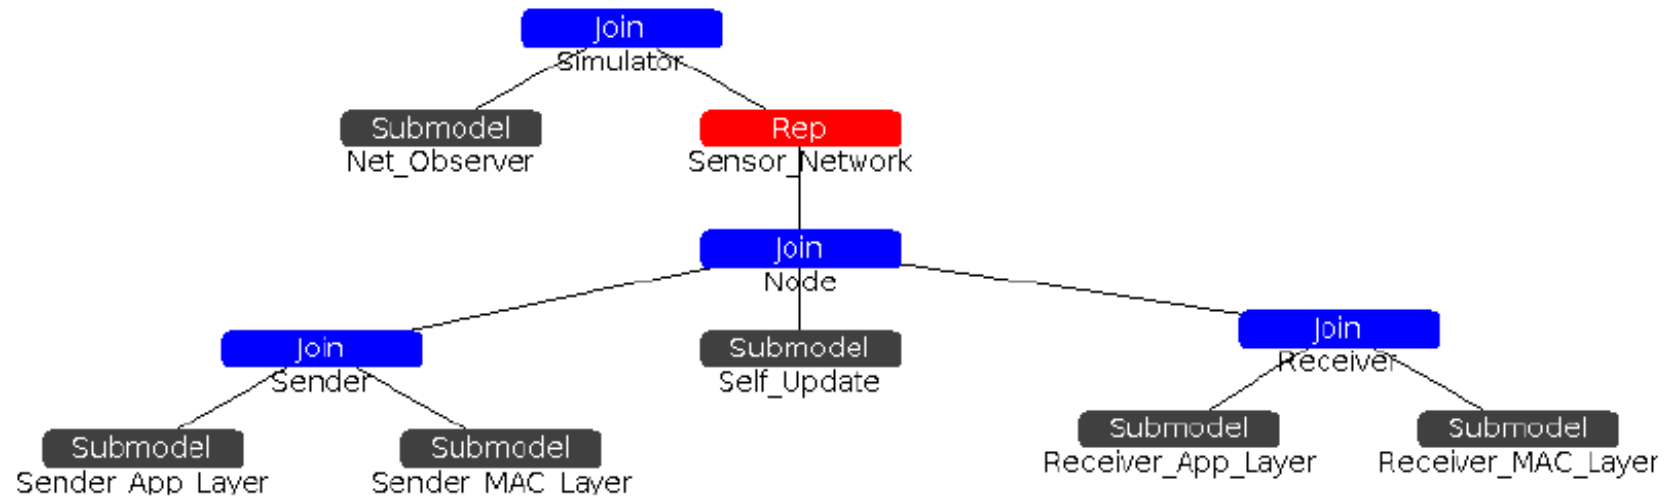
Mobius tool provides a comprehensive framework for evaluation of system dependability and performance.

The main features of the tool include:

- (i) a hierarchical modeling paradigm, allowing one to build complex models by specifying the behavior of individual components and combining them to create a complete system model;
- (ii) customised measures of system properties;
- (iii) distributed discrete-event simulation;
- (iv) numerical solution of Markov models.

Hierarchical model: an example

sensor network protocol + network observer



Rep operator:

constructs a number of replicas of an individual component

Join operator:

groups individual and/or replicated components.

Activities: sensor network protocol

Activities

Activity	Duration	Description
Sender MAC Layer		
start	instantaneous	transmission start
test	instantaneous	listening the channel
backoff	exponential	backoff period
tx_chunk	$c(\text{chunk_length} \rightarrow \text{Mark}())$	transmit the preamble
tx_data	constant	transmit the data packet
ready_to_send	instantaneous	check interval end
Receiver MAC Layer		
duty_on	constant	wake up and channel check
rx_chunk	exponential	reception of chunk
sleep	$c'(\text{sleep_length} \rightarrow \text{Mark}())$	sleep period between chunk and data packet
rx_data	constant	reception of a data packet
duty_off	$c''(\text{duty_off_length} \rightarrow \text{Mark}())$	duty off period
restart	instantaneous	check interval end/restart
Self_Update		
time	constant	periodic update of timers
update	instantaneous	update timers

Simulation and Reward functions

Mobius allows system simulation and allows properties of interest to be specified with **reward functions**

ASSIGNING REWARDS TO STATES or TRANSITIONS defines a Markov reward model.


RATE rewards (in the case of assignment to states)

or

IMPLULSE rewards (in the case of assignment to transition)

Simulation and Reward functions

RATE rewards



Assume place A in the model is the only state that represents the failure of the system

Consider the performance variable Rel, which measures reliability

Generally the value of the performance variable is the **mean** of the values returned by the associated reward function

The reward function is:

```
if (Model-name->A->Mark()==0) return 1;  
    else return 0;
```

Simulation and Reward functions

Reward function:

specifies how to measure a property on the basis of the SAN marking.

Measurements can be made at specific time instants, over periods of time, or when the system reaches a steady state.

A desired confidence level is associated to each reward function.

At the end of a simulation the Mobius tool is able to evaluate for each reward function whether the desired confidence level has been attained or not thus ensuring high accuracy of measurements.

Confidence interval: accuracy of the measure of the property

Solutions on measures

Numerical solver

model is the reachability graph.

Only for models with deterministic and exponentially distributed activities

Finite state, small states-space description

Simulation

model is a discrete event simulator.

For all models

- arbitrary large state-space description
- provide statistically accurate solutions within some user specifiable confidence interval